




Accuracy (%), seeds averaged per subject. Rows ordered by mean over corpora. Each corpus has its OWN colour scale (they span different ranges), so compare within a panel, not across panels. A dash is a model that cannot run on that corpus.